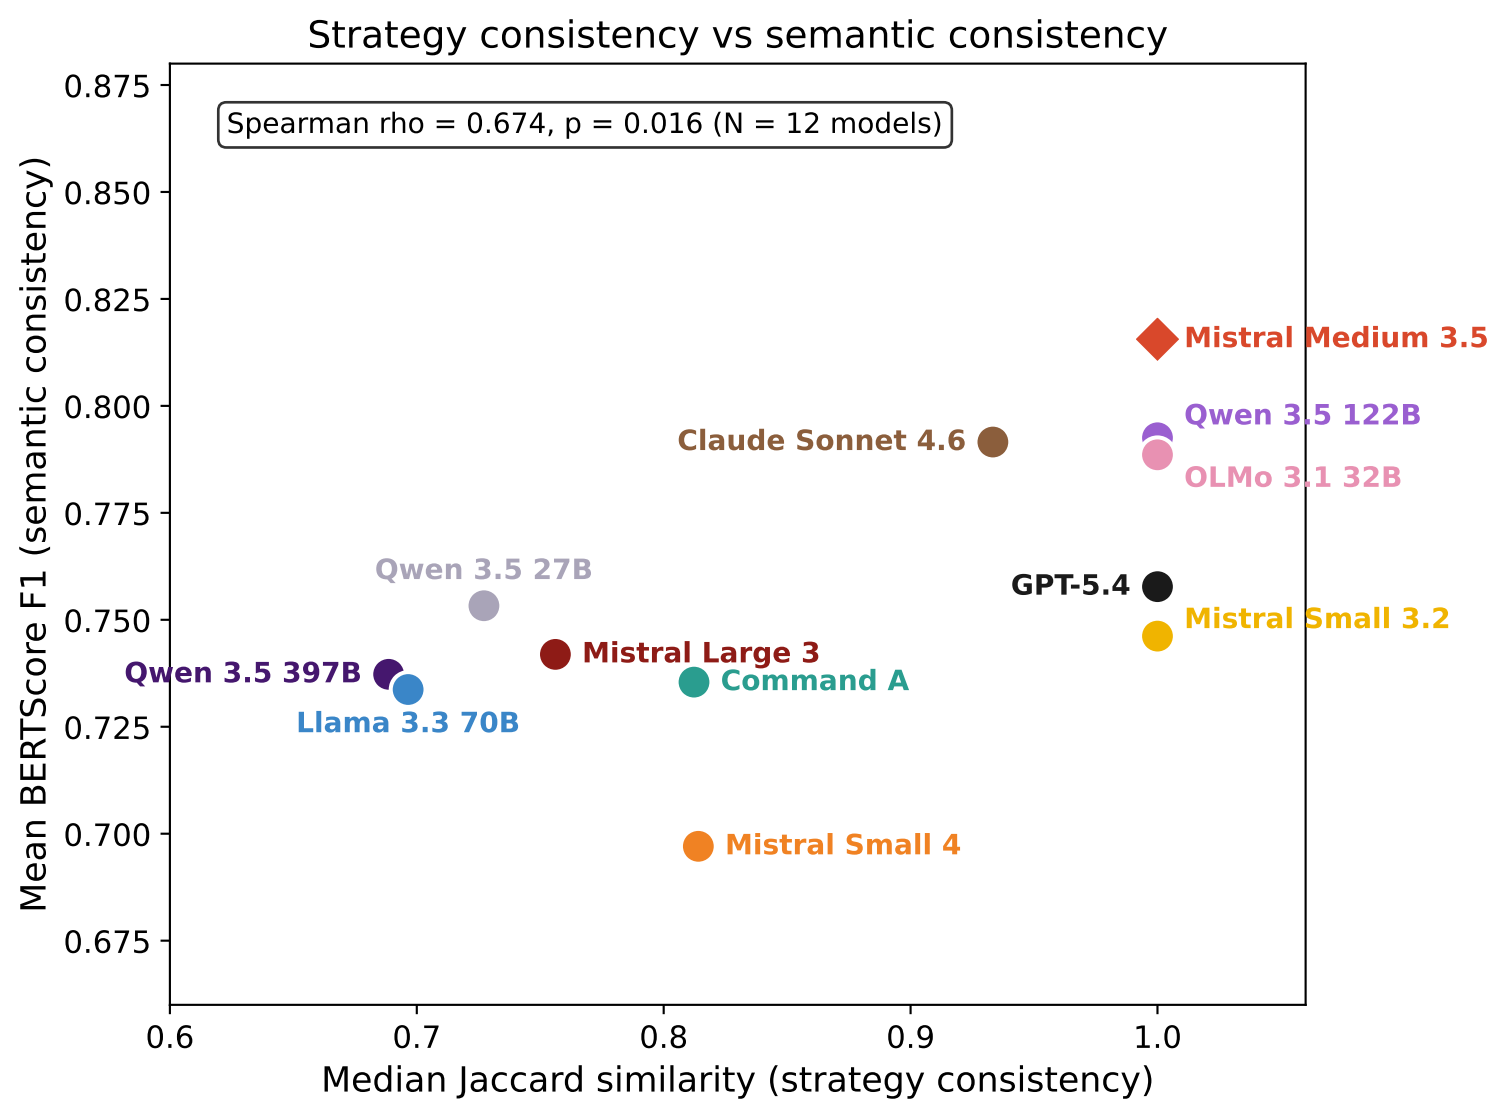
Strategy consistency vs semantic consistency
Spearman rho = 0.674, p = 0.016 (N = 12 models)
Mean BERTScore F1 (semantic consistency)
Median Jaccard similarity (strategy consistency)
0.875
0.850
0.825
0.800
0.775
0.750
0.725
0.700
0.675
0.6
0.7
0.8
0.9
1.0
Mistral Medium 3.5
Qwen 3.5 122B
Claude Sonnet 4.6
OLMo 3.1 32B
Qwen 3.5 27B
GPT-5.4
Mistral Small 3.2
Mistral Large 3
Qwen 3.5 397B
Command A
Llama 3.3 70B
Mistral Small 4